
Network Security Basics

Outline

- **IP Address and Network Interface**
- **TCP/IP Protocols**
- **Packet Sniffing**
- **Packet Spoofing**
- **Programming using Scapy**
- **Lab environment and containers**

IP ADDRESS

IP Address: the Original Scheme

```

Class A      |<- -      Host ID      -->|
  0.   0.   0.   0 = 00000000.00000000.00000000.00000000
127.255.255.255 = 01111111.11111111.11111111.11111111

```

```

Class B                                     |<-- Host ID -->|
128.  0.  0.  0 = 100000000.000000000.000000000.000000000
191.255.255.255 = 101111111.111111111.111111111.111111111

```

```

Class C                                     |HostID|
192.  0.  0.  0 = 11000000.00000000.00000000.00000000
223.255.255.255 = 11011111.11111111.11111111.11111111

```

```

Class D           |<--      Address Range      -->|
224.  0.  0.  0 = 11100000.00000000.00000000.00000000
239.255.255.255 = 11101111.11111111.11111111.11111111

```

```

Class E                |<--      Address Range      -->|
240.  0.  0.  0 = 11110000.00000000.00000000.00000000
255.255.255.255 = 11111111.11111111.11111111.11111111

```

CIDR Scheme (Classless Inter-Domain Routing)

192.168.60.5/24



Indicate the first
24 bits are
network ID

Question: What is the address range of the network **192.168.192.0/19** ?

Special IP Addresses

- **Private IP Addresses**

- 10.0.0.0/8
- 172.16.0.0/12
- 192.168.0.0/16

- **Loopback Address**

- 127.0.0.0/8
- Commonly used: 127.0.0.1

List IP Address on Network Interface

```
$ ip -br address
lo                UNKNOWN    127.0.0.1/8 ::1/128
enp0s3           UP          10.0.5.5/24 fe80::bed8:53e2:5192:f265/64
docker0          DOWN        172.17.0.1/16 fe80::42:13ff:fee7:90d6/64
```

Manually Assign IP Address

```
$ sudo ip addr add 192.168.60.6/24 dev enp0s3
$ ip addr
1: lo: <LOOPBACK,UP,LOWER_UP> mtu 65536 qdisc noqueue state UNKNOWN
    group default qlen 1
    link/loopback 00:00:00:00:00:00 brd 00:00:00:00:00:00
    inet 127.0.0.1/8 scope host lo
        valid_lft forever preferred_lft forever
    inet6 ::1/128 scope host
        valid_lft forever preferred_lft forever
2: enp0s3: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc pfifo_f
ast state UP group default qlen 1000
    link/ether 08:00:27:84:5e:b9 brd ff:ff:ff:ff:ff:ff
    inet 192.168.60.6/24 scope global enp0s3
        valid_lft forever preferred_lft forever
    inet6 fe80::3fc4:1dac:bbbb:948/64 scope link
        valid_lft forever preferred_lft forever
```


Automatically Assign IP Address

- **DHCP: Dynamic Host Configuration Protocol**

Get IP Addresses for Host Names: DNS

```
seed@VM:~$ dig www.example.com
```


```
; <<>> DiG 9.16.1-Ubuntu <<>> www.example.com
;; global options: +cmd
;; Got answer:
;; ->>HEADER<<- opcode: QUERY, status: NOERROR, id: 18093
;; flags: qr rd ra; QUERY: 1, ANSWER: 1, AUTHORITY: 0, ADDITIONAL: 1

;; OPT PSEUDOSECTION:
; EDNS: version: 0, flags::; udp: 65494
;; QUESTION SECTION:
;www.example.com.                IN      A

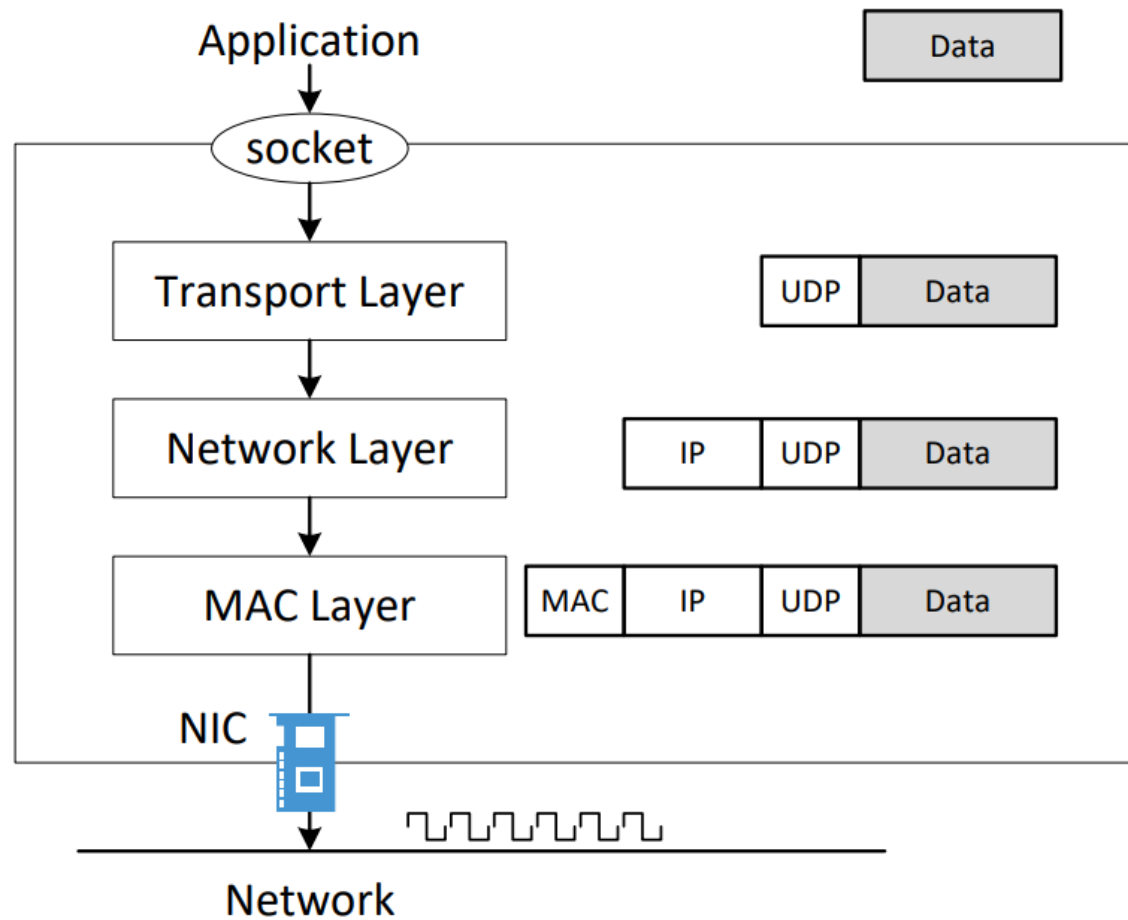
;; ANSWER SECTION:
www.example.com.                57405   IN      A      93.184.216.34
```

NETWORK STACK

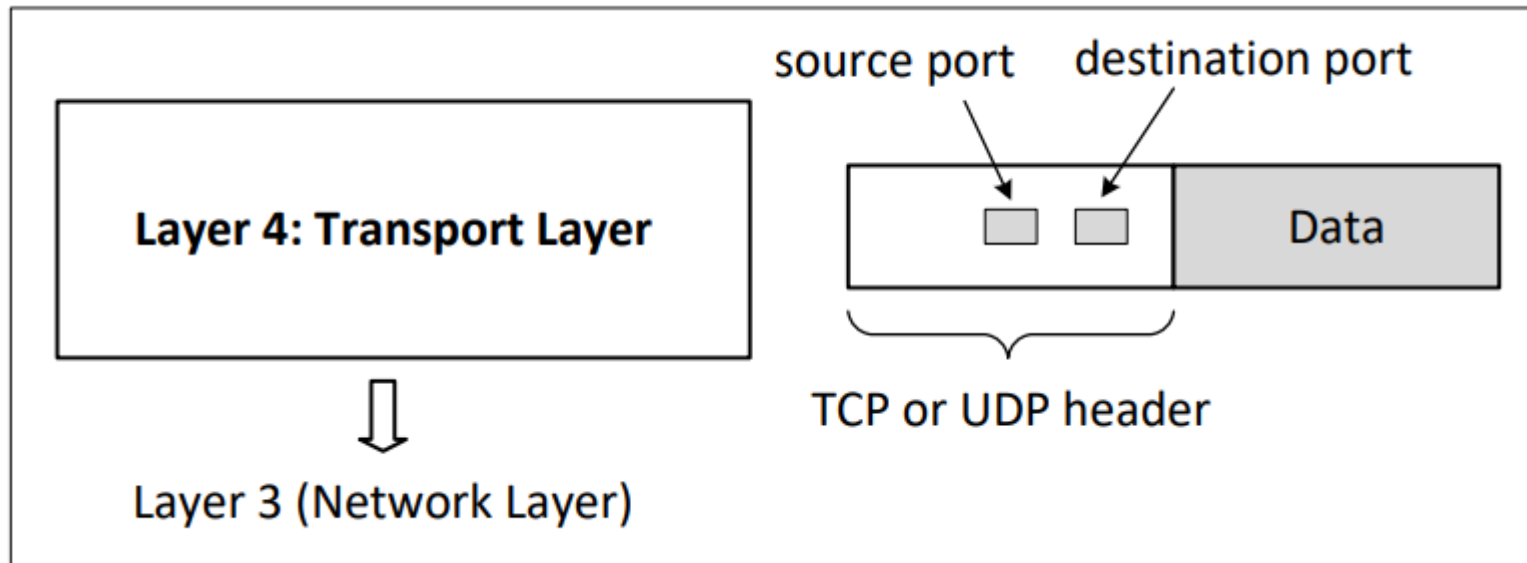
Packet Journey at High Level



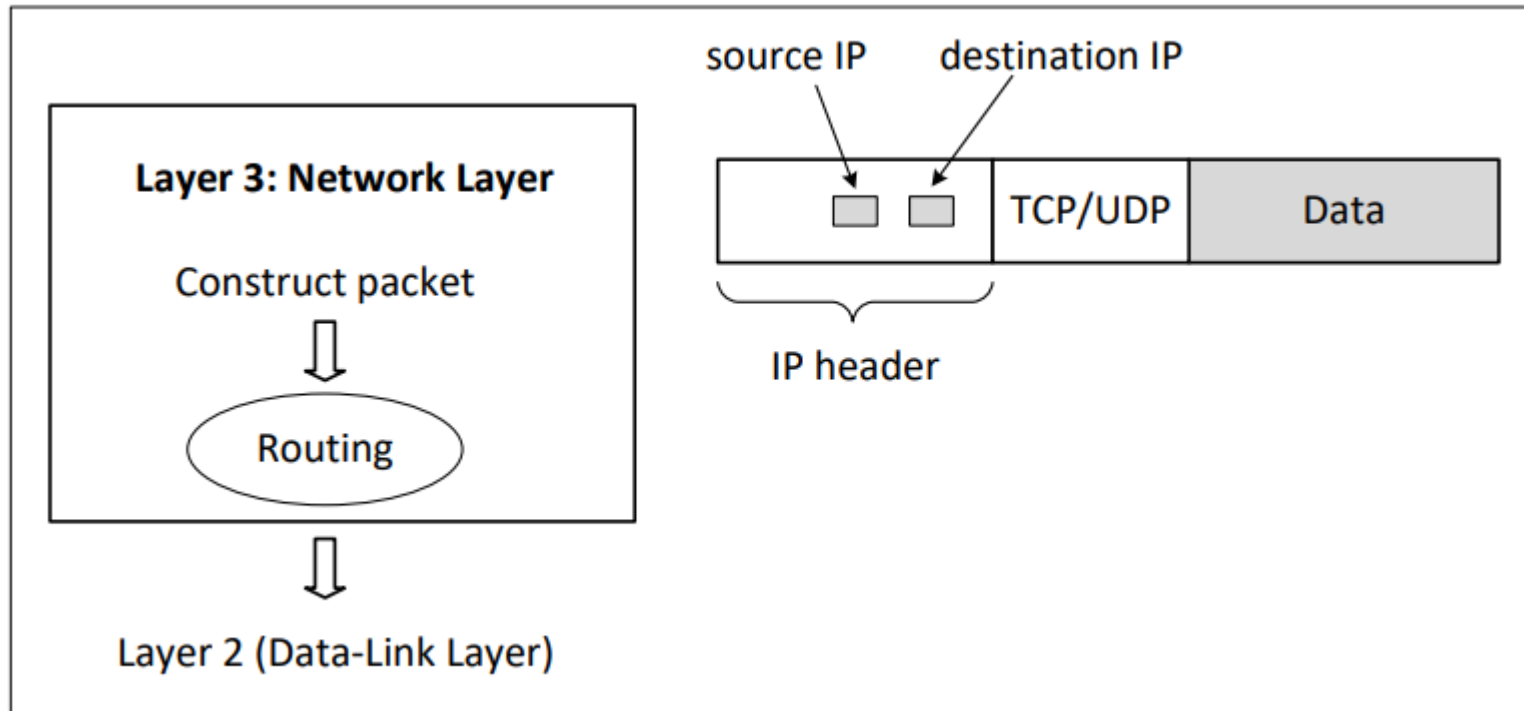
How Packets Are Constructed



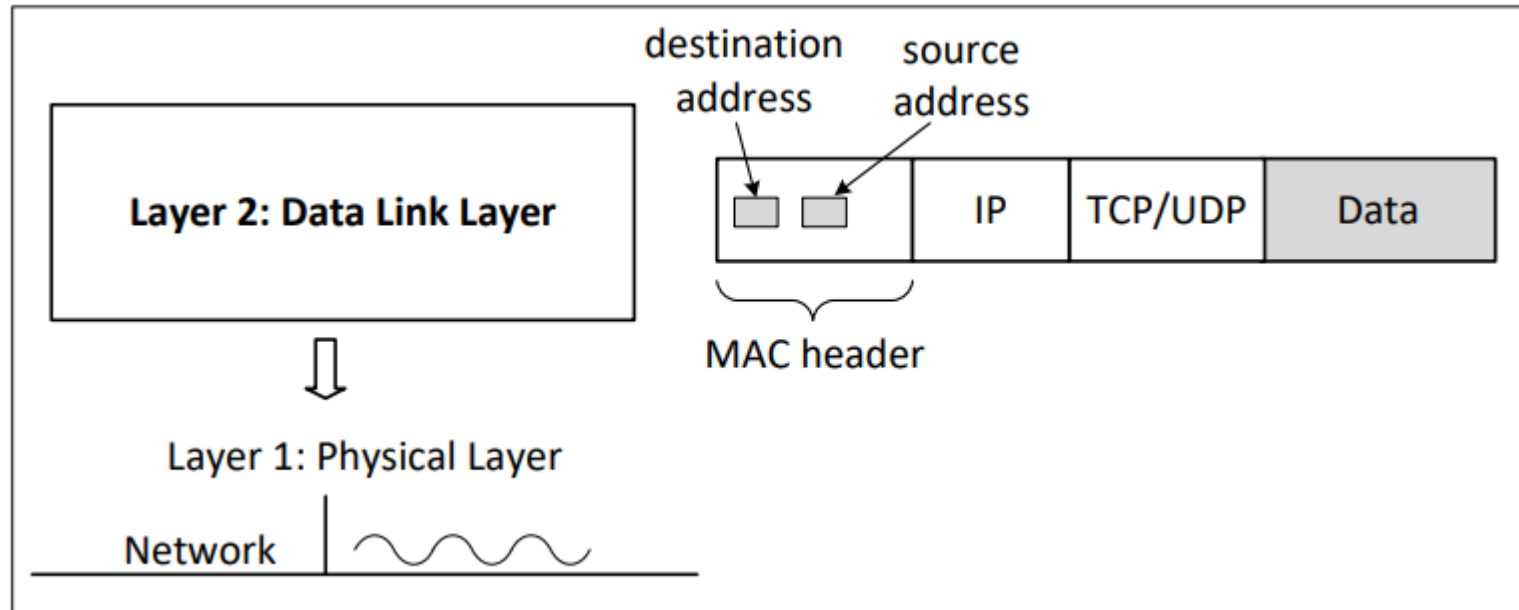
Layer 4: Transport Layer



Layer 3: Network Layer



Layer 2: Data Link Layer (MAC Layer)



Sending Packet in Python

- **UDP Client**

```
#!/usr/bin/python3
```

```
import socket
```

```
IP    = "127.0.0.1"
```

```
PORT = 9090
```

```
data = b'Hello, World!'
```

```
sock = socket.socket(socket.AF_INET, socket.SOCK_DGRAM)
```

```
sock.sendto(data, (IP, PORT))
```

Sending Packet in Python (1)

- Execution Results

```
$ nc -l uv 9090
Listening on [0.0.0.0] (family 0, port 9090)
Hello, World!█
```

Receiving Packets in Python

- **UDP Server**

```
#!/usr/bin/python3

import socket

IP    = "0.0.0.0"
PORT = 9090

sock = socket.socket(socket.AF_INET, socket.SOCK_DGRAM)
sock.bind((IP, PORT))

while True:
    data, (ip, port) = sock.recvfrom(1024)
    print("Sender: {} and Port: {}".format(ip, port))
    print("Received message: {}".format(data))
```

UDP Server

Terminal

```
seed@10.0.2.6:$ nc -u 10.0.2.7 9090
```

```
hello
```

```
hello again
```

```
█
```

Terminal

```
Server(10.0.2.7):$ udp_server.py
```

```
Sender: 10.0.2.6 and Port: 49112
```

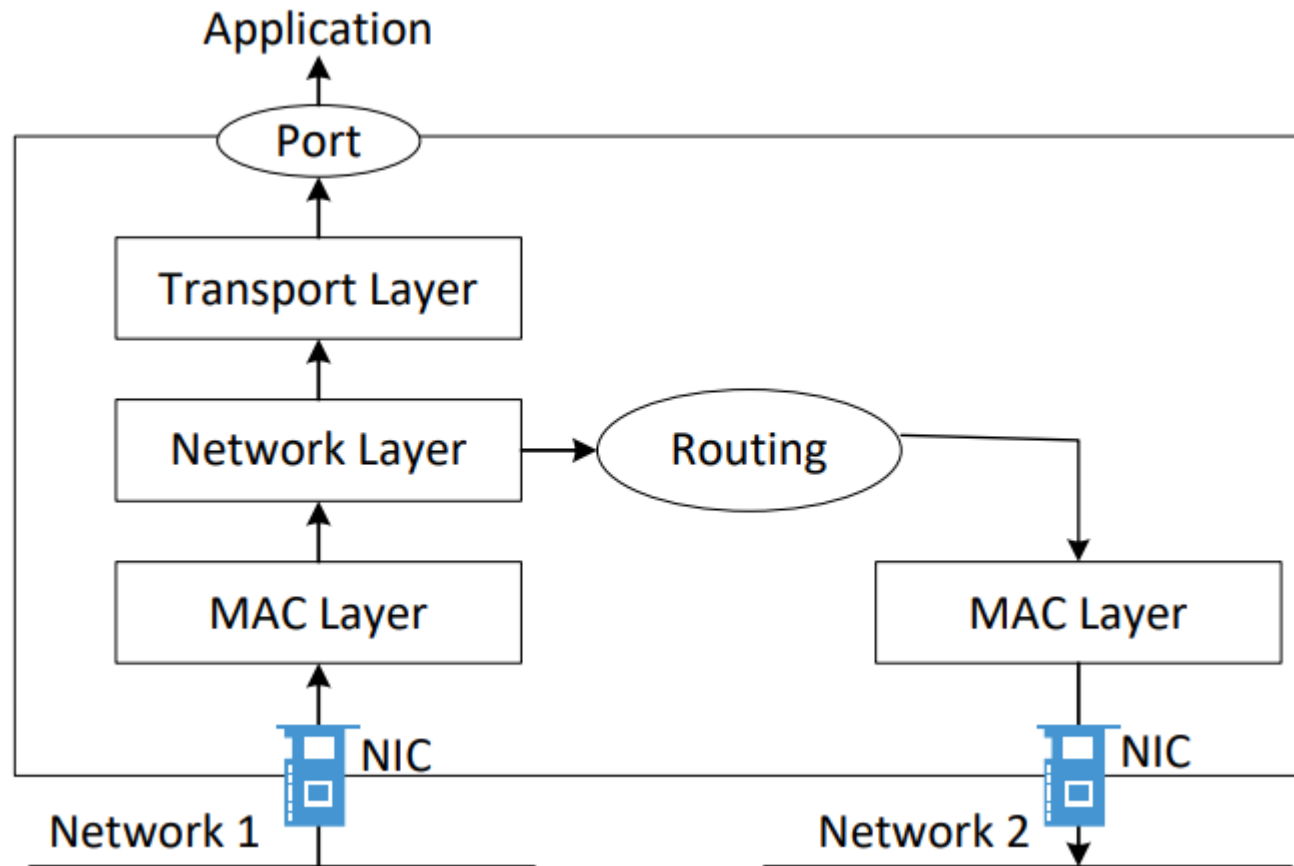
```
Received message: b'hello\n'
```

```
Sender: 10.0.2.6 and Port: 49112
```


```
Received message: b'hello again\n'
```

```
█
```

How Packets Are Received



Routing



The “ip route” Command

```
# ip route
default via 10.9.0.1 dev eth0
10.9.0.0/24 dev eth0 proto kernel scope link src 10.9.0.11
192.168.60.0/24 dev eth1 proto kernel scope link src 192.168.60.11

# ip route get 10.9.0.1
10.9.0.1 dev eth0 src 10.9.0.11 uid 0

# ip route get 192.168.60.5
192.168.60.5 dev eth1 src 192.168.60.11 uid 0

# ip route get 1.2.3.4
1.2.3.4 via 10.9.0.1 dev eth0 src 10.9.0.11 uid 0
```

Packet Sending Tools

- Using netcat

```
$ nc <ip> <port>      ← send out TCP packet  
$ nc -u <ip> <port>   ← send out UDP packet
```

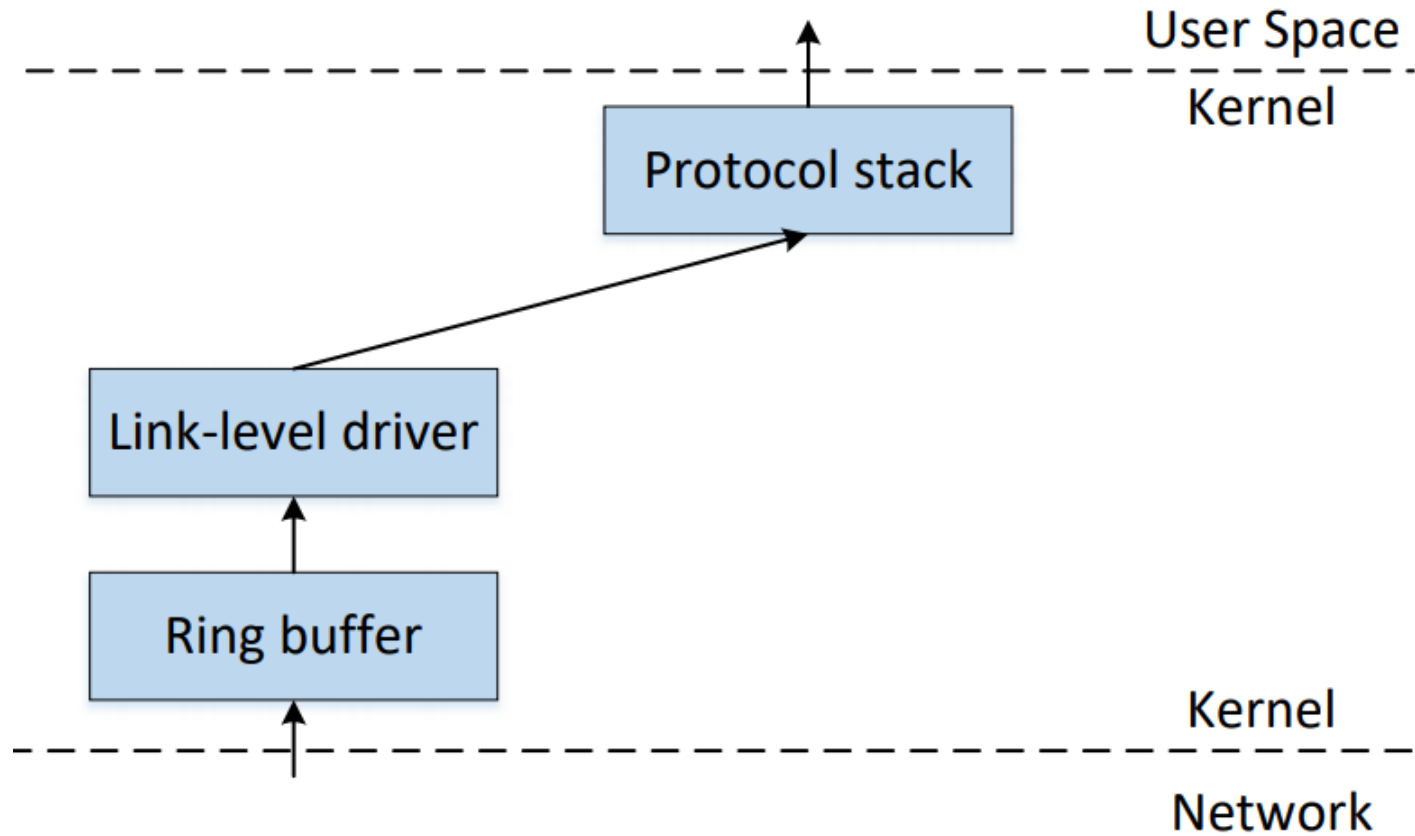
- Bash: /dev/tcp or /dev/udp pseudo device

```
$ echo "data" > /dev/udp/<ip>/<port>  
$ echo "data" > /dev/tcp/<ip>/<port>
```

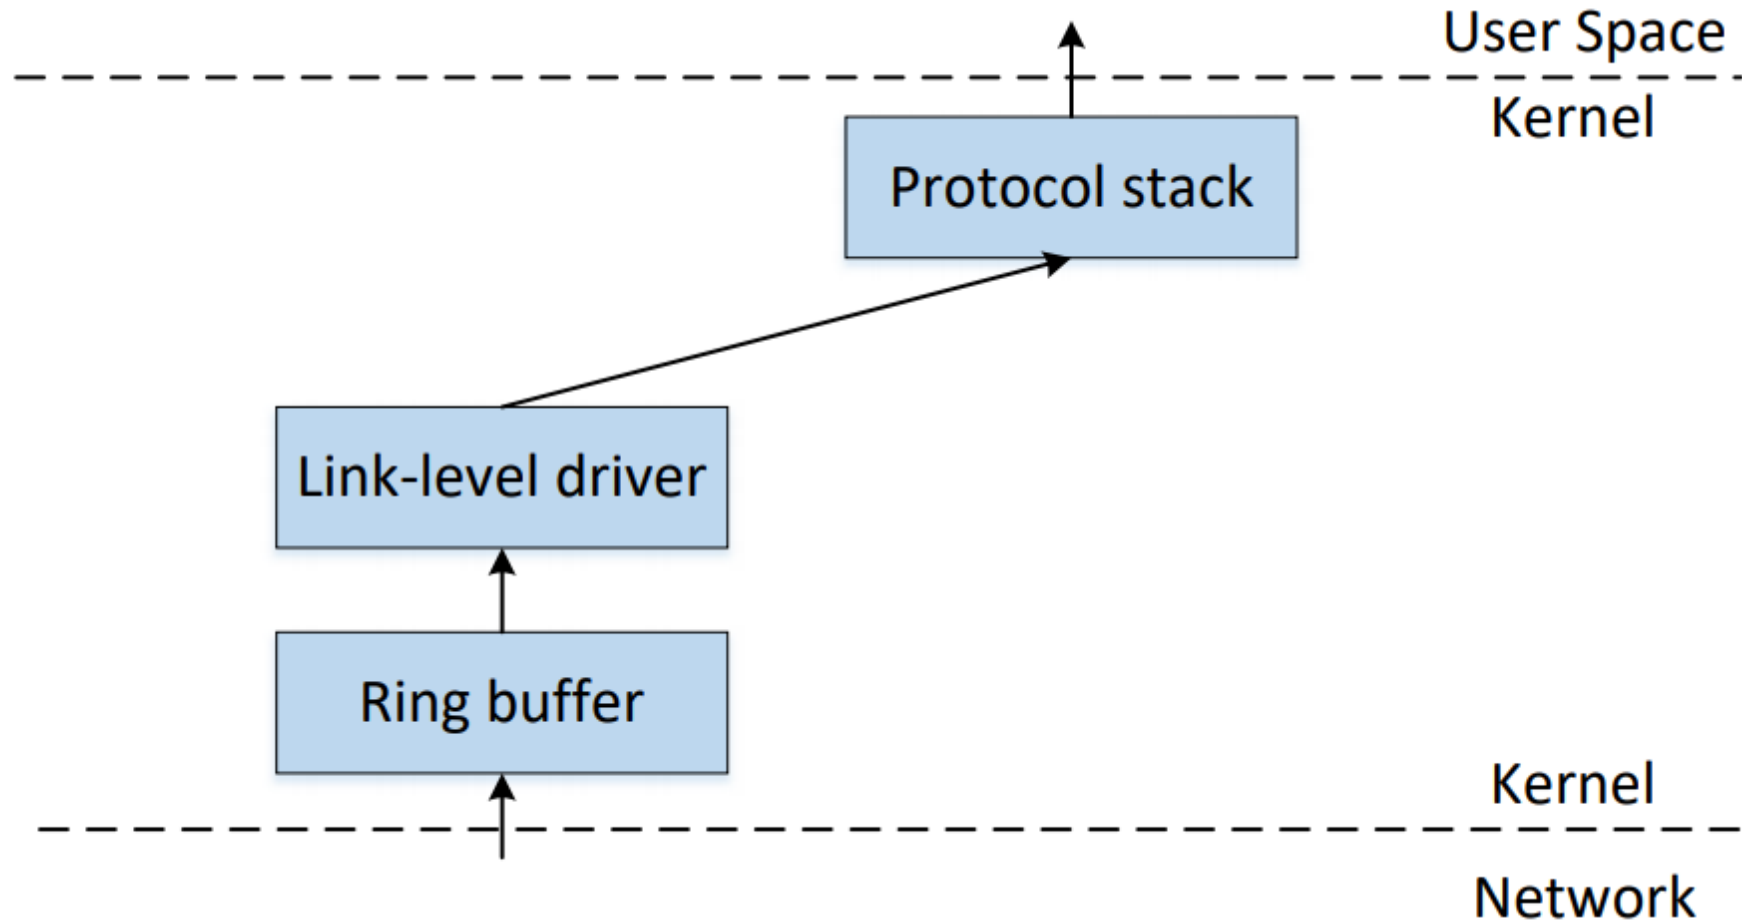
- Others: telnet, ping, etc.

PACKET SNIFFING


How Packets Are Received



How To Get A Copy of Packet



Lab Setup



```
seed@VM:~$ dockps
9eb2c057887f  host-10.9.0.5
89a0dfac1c75  host-10.9.0.6
f452376e85a5  host-192.168.60.5
8856896b15ea  host-192.168.60.6
9aa28fadb047  router
```

Packet Sniffing Tools

- **Tcpdump**

- Command line
- Good choice for containers (in the lab setup)

- **Wireshark**

- GUI
- Good choices for the environment supporting GUI (not containers)

- **Scapy**


- Implement your own sniffing tools

Tcpdump Examples

- `tcpdump -n -i eth0`
 - `-n`: do not resolve the IP address to host name
 - `-i`: sniffing on this interface
- `tcpdump -n -i eth0 -vvv "tcp port 179"`
 - `-vvv`: asks the program to produce more verbose output.
- `tcpdump -i eth0 -w /tmp/packets.pcap`
 - saves the captured packets to a PCAP file
 - use Wireshark to display them

Wireshark and Containers

Find the correct interface



```
seed@VM:~$ docker network ls
NETWORK ID          NAME                DRIVER              SCOPE
d10f14b6b6f9        bridge              bridge              local
b3581338a28d        host                host                local
54cdc1101b1a        net-10.9.0.0        bridge              local
0569b491802f        net-192.168.60.0    bridge              local
77aceccccbe26        none                null                local

seed@VM:~$ ip -br address
lo                UNKNOWN           127.0.0.1/8 ::1/128
enp0s3            UP                10.0.5.5/24 fe80::bed8:53e2:5192:f265/64
docker0           DOWN              172.17.0.1/16 fe80::42:13ff:fee7:90d6/64
br-54cdc1101b1a   UP                10.9.0.1/24 fe80::42:1cff:fe17:f3e6/64
br-0569b491802f   UP                192.168.60.1/24 fe80::42:b5ff:fe9b:6b49/64
```

Scapy Example 1

```
#!/usr/bin/python3

from scapy.all import *

pkt = sniff(iface='enp0s3',
            filter='icmp or udp',
            count=10)

pkt.summary()
```

```
seed@VM:~$ ip -br addr
lo                UNKNOWN      127.0.0.1/8 ::1/
enp0s3            UP           10.0.5.5/24 fe80:
docker0           DOWN        172.17.0.1/16 fe
br-54cdc1101b1a   UP           10.9.0.1/24 fe80:
br-0569b491802f   UP           192.168.60.1/24
```

```
root@9eb2c057887f:~# ip -br addr
lo                UNKNOWN      127.0.0.1/8
eth0@if1882       UP           10.9.0.5/24
```


Scapy Example 2

```
#!/usr/bin/python3

from scapy.all import *

def process_packet(pkt):
    #hexdump(pkt)
    pkt.show()
    print("-----")

f = 'udp and dst portrange 50-55 or icmp'

sniff(iface='enp0s3', filter = f, prn=process_packet)
```

Filter Examples for Scapy

- Berkeley Packet Filter (BPF) syntax
- Same as tcpdump

```
dst host 10.0.2.5: only capture the packets going to 10.0.2.5.  
src host 10.0.2.6: only capture the packets coming from 10.0.2.6.  
host 10.0.2.6 and src port 9090: only capture the packets coming  
from or going to 10.0.2.6 with the source port being 9090.  
tcp: only capture TCP packets.
```

Scapy: Display Packets

- Using `hexdump()`

```
>>> hexdump(pkt)
0000  52 54 00 12 35 00 08
0010  00 54 F2 29 40 00 40
0020  08 08 08 00 98 01 10
0030  0C 00 08 09 0A 0B 0C
0040  16 17 18 19 1A 1B 1C
0050  26 27 28 29 2A 2B 2C
0060  36 37
```

- Using `pkt.show()`

```
>>> pkt.show()
###[ Ethernet ]###
    dst      = 52:54:00:12:35:00
    src      = 08:00:27:77:2e:c3
    type     = IPv4
###[ IP ]###
    version  = 4
    ihl      = 5
    ...
    proto    = icmp
    chksum   = 0x3c9a
    src      = 10.0.2.8
    dst      = 8.8.8.8
    \options \
###[ ICMP ]###
```

Scapy: Iterate Through Layers

```
>>> pkt = Ether()/IP()/UDP()/ "hello"  
>>> pkt  
<Ether type=IPv4 |<IP frag=0 proto=udp |<UDP |<Raw load='hello' |>>>>
```

```
>>> pkt.payload                                ← an IP object  
<IP frag=0 proto=udp |<UDP |<Raw load='hello' |>>>  
  
>>> pkt.payload.payload                        ← a UDP object  
<UDP |<Raw load='hello' |>>  
  
>>> pkt.payload.payload.payload                ← a Raw object  
<Raw load='hello' |>  
  
>>> pkt.payload.payload.payload.load           ← the actual payload  
b'hello'
```

Accessing Layers

Get inner layers

```
>>> pkt.getlayer(UDP)
<UDP  |<Raw  load='hello'  |>>
>>> pkt[UDP]
<UDP  |<Raw  load='hello'  |>>

>>> pkt.getlayer(Raw)
<Raw  load='hello'  |>
>>> pkt[Raw]
<Raw  load='hello'  |>
```

Check layer existence

```
>>> pkt.haslayer(UDP)
True
>>> pkt.haslayer(TCP)
0
>>> pkt.haslayer(Raw)
True
```

A Sniffer Example

```
def process_packet(pkt):
    if pkt.haslayer(IP):
        ip = pkt[IP]
        print("IP: {} --> {}".format(ip.src, ip.dst))

    if pkt.haslayer(TCP):
        tcp = pkt[TCP]
        print("    TCP  port: {} --> {}".format(tcp.sport, tcp.dport))

    elif pkt.haslayer(UDP):
        udp = pkt[UDP]
        print("    UDP  port: {} --> {}".format(udp.sport, udp.dport))

    elif pkt.haslayer(ICMP):
        icmp = pkt[ICMP]
        print("    ICMP type: {}".format(icmp.type))

    else:
        print("    Other protocol")

sniff(iface='enp0s3', filter='ip', prn=process_packet)
```

PACKET SPOOFING

Packet Spoofing

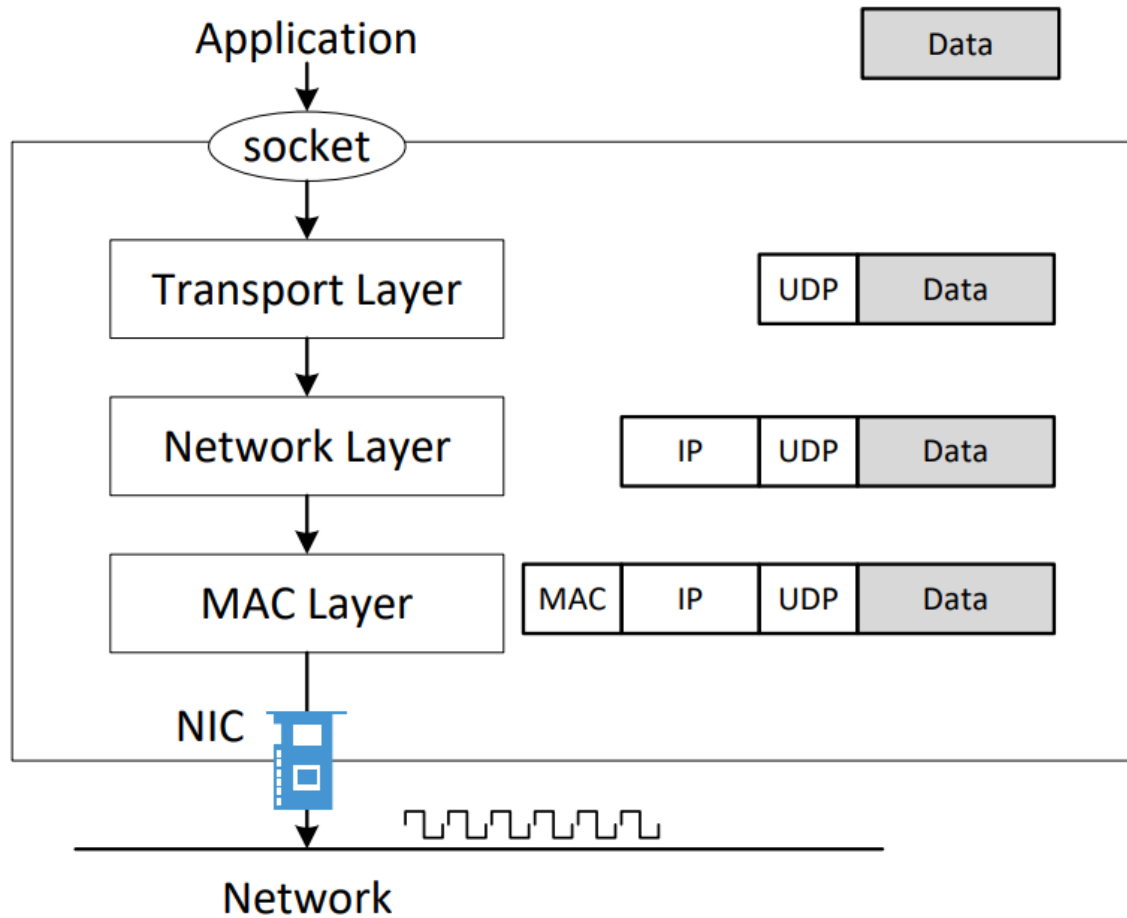
- **In normal packet construction**

- Only some selected header fields can be set by users
- OS set the other fields

- **Packet spoofing**

- Set arbitrary header fields
- Using tools
- Using Scapy

How To Spoof Packets



Spoofing ICMP Packets

```
#!/usr/bin/python3
from scapy.all import *

print("SENDING SPOOFED ICMP PACKET.....")
ip = IP(src="1.2.3.4", dst="93.184.216.34")
icmp = ICMP()
pkt = ip/icmp
pkt.show()
send(pkt, verbose=0)
```

Spoofing UDP Packets

```
#!/usr/bin/python3
from scapy.all import *

print("SENDING SPOOFED UDP PACKET.....")
ip = IP(src="1.2.3.4", dst="10.0.2.69") # IP Layer
udp = UDP(sport=8888, dport=9090)      # UDP Layer
data = "Hello UDP!\n"                  # Payload
pkt = ip/udp/data
pkt.show()
send(pkt, verbose=0)
```

Sniff Request and Spoof Reply: Code

```
def spoof_pkt(pkt):
    if ICMP in pkt and pkt[ICMP].type == 8:
        print("Original Packet.....")
        print("Source IP : ", pkt[IP].src)
        print("Destination IP :", pkt[IP].dst)

        ip = IP(src=pkt[IP].dst, dst=pkt[IP].src,
                ihl=pkt[IP].ihl, ttl = 99)
        icmp = ICMP(type=0, id=pkt[ICMP].id, seq=pkt[ICMP].seq)
        data = pkt[Raw].load
        newpkt = ip/icmp/data

        print("Spoofed Packet.....")
        print("Source IP : ", newpkt[IP].src)
        print("Destination IP :", newpkt[IP].dst)

        send(newpkt, verbose=0)

pkt = sniff(iface = 'br-54cdc1101b1a',
            filter = 'icmp and src host 10.9.0.5',
            prn = spoof_pkt)
```

Other Uses of Scapy: Send and Receive

- `send()` : Send packets at Layer 3.
- `sendp()` : Send packets at Layer 2.
- `sr()` : Sends packets at Layer 3 and receiving answers.
- `srp()` : Sends packets at Layer 2 and receiving answers.
- `sr1()` : Sends packets at Layer 3 and waits for the first answer.
- `sr1p()` : Sends packets at Layer 2 and waits for the first answer.
- `srloop()` : Send a packet at Layer 3 in a loop and print the answer each time.
- `srploop()` : Send a packet at Layer 2 in a loop and print the answer each time.

Example: implement ping

```
#!/usr/bin/python3
from scapy.all import *

ip = IP(dst="8.8.8.8")
icmp = ICMP()
pkt = ip/icmp
reply = sr1(pkt)
print("ICMP reply .....")
print("Source IP : ", reply[IP].src)
print("Destination IP :", reply[IP].dst)
```

Traceroute Code

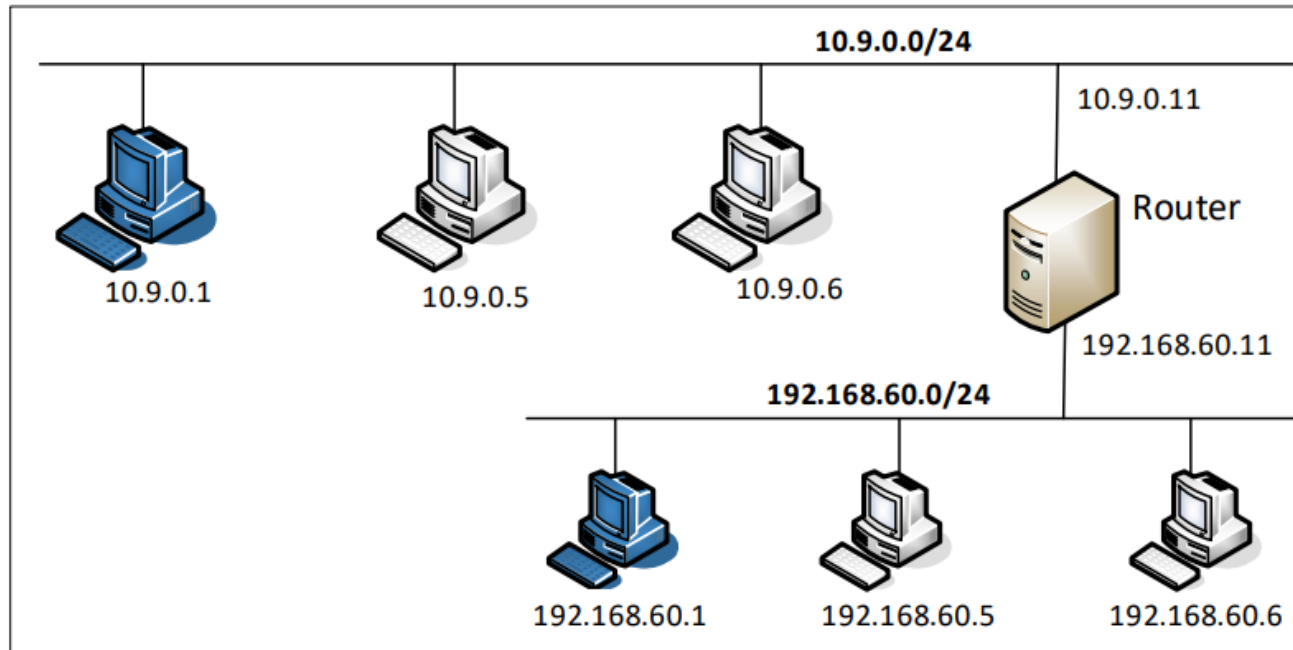
```
b = ICMP()  
a = IP()  
a.dst = '93.184.216.34'  
  
TTL = 3  
a.ttl = TTL  
h = srl(a/b, timeout=2, verbose=0)  
if h is None:  
    print("Router: *** (hops = {})".format(TTL))  
else:  
    print("Router: {} (hops = {})".format(h.src, TTL))
```

DOCKER SETUP

Lab Setup and Containers

- **Most labs in Internet Security use containers**

- **Lab setup files::** [Labsetup.zip](#)
- **Manual::** [Docker manual](#)



Docker Compose

- **Setup file:** `docker-compose.yml`

```
version: "3"

services:
  HostA1:
    ...
  HostA2:
    ...
  ...
networks:
  net-192.168.60.0:
    ...
  net-10.9.0.0:
    ...
```

Docker Manual:
<https://github.com/seed-labs/seed-labs/blob/master/manuals/docker/SEEDManual-Container.md>

Set Up Networks

```
networks:
  net-10.9.0.0:
    name: net-10.9.0.0
    ipam:
      config:
        - subnet: 10.9.0.0/24

  net-192.168.60.0:
    name: net-192.168.60.0
    ipam:
      config:
        - subnet: 192.168.60.0/24
```

Find out interface name

```
$ ifconfig
br-03bc5aebc4c4: flags=4163<UP,
                  inet 10.9.0.1 netmask
```

```
$ docker network ls
NETWORK ID          NAME
c616fa7f4f46        bridge
b3581338a28d        host
03bc5aebc4c4        net-10.9.0.0
e0afdc1c0e70        net-192.168.60.0
```

Set Up Hosts

```
HostA1:
  image: handsonsecurity/seed-ubuntu:large
  container_name: host-10.9.0.5
  tty: true
  cap_add:
    - ALL
  privileged: true
  volumes:
    - ./volumes:/volumes
  networks:
    net-10.9.0.0:
      ipv4_address: 10.9.0.5
  command: bash -c "
              ip route add 192.168.60.0/24 via 10.9.0.11 &&
              tail -f /dev/null
            "
```

Sniffing Inside Containers

- **Limitation**

- Can only sniff its own traffic
- Due to how the virtual network is implemented



Sniffing Inside Containers

- **Overcome the limitation**
 - Use the “host” mode
 - `network_mode: host`

Start/Stop Containers

Alias created in the SEED VM

```
docker-compose build  
docker-compose up  
docker-compose down
```

```
dcbuild  
dcup  
dcdow
```

Get Into A Container

Alias created in the SEED VM

```
$ docker ps
CONTAINER ID        NAMES                ...
bcff498d0b1f        host-10.9.0.6        ...
1e122cd314c7        host-10.9.0.5        ...
31bd91496f62        host-10.9.0.7        ...

$ docker exec -it 1e /bin/bash
root@1e122cd314c7:/#
```

```
$ dockps
bcff498d0b1f  host-10.9.0.6
1e122cd314c7  host-10.9.0.5
31bd91496f62  host-10.9.0.7

$ docksh 31
root@31bd91496f62:/#
```


Copy Files Between Host and Container

Get container ID

```
$ docker ps
CONTAINER ID        NAMES
bcff498d0b1f        host-10.9.0.6
1e122cd314c7        host-10.9.0.5
31bd91496f62        host-10.9.0.7
```

```
// From host to container
$ docker cp file.txt bcff:/tmp/
$ docker cp folder bcff:/tmp

// From container to host
$ docker cp bcff:/tmp/file.txt .
$ docker cp bcff:/tmp/folder .
```